Report

1. Task 1 (pruning.py)
   1.1 Model summary

```
Model: "model"
_________________________________________________________________________________________________
Layer (type)                 Output Shape              Param #   Connected to
=================================================================================================
input_1 (InputLayer)         [(None, 96, 96, 3)]       0
_________________________________________________________________________________________________
prune_low_magnitude_Conv1_pad ( (None, 97, 97, 3)      1         input_1[0][0]
_________________________________________________________________________________________________
prune_low_magnitude_Conv1 (Prun (None, 48, 48, 32)     1730      prune_low_magnitude_Conv1_pad[0][
_________________________________________________________________________________________________
prune_low_magnitude_bn_Conv1 (P (None, 48, 48, 32)     129       prune_low_magnitude_Conv1[0][0]
_________________________________________________________________________________________________
prune_low_magnitude_Conv1_relu  (None, 48, 48, 32)     1         prune_low_magnitude_bn_Conv1[0][0
_________________________________________________________________________________________________
prune_low_magnitude_expanded_co (None, 48, 48, 32)     289       prune_low_magnitude_Conv1_relu[0]
_________________________________________________________________________________________________
prune_low_magnitude_expanded_co (None, 48, 48, 32)     129       prune_low_magnitude_expanded_conv
_________________________________________________________________________________________________
prune_low_magnitude_expanded_co (None, 48, 48, 32)     1         prune_low_magnitude_expanded_conv
_________________________________________________________________________________________________
prune_low_magnitude_expanded_co (None, 48, 48, 16)     1026      prune_low_magnitude_expanded_conv
_________________________________________________________________________________________________
prune_low_magnitude_expanded_co (None, 48, 48, 16)     65        prune_low_magnitude_expanded_conv
_________________________________________________________________________________________________
prune_low_magnitude_block_1_exp (None, 48, 48, 96)     3074      prune_low_magnitude_expanded_conv
_________________________________________________________________________________________________
prune_low_magnitude_block_1_exp (None, 48, 48, 96)     385       prune_low_magnitude_block_1_expan
_________________________________________________________________________________________________
prune_low_magnitude_block_1_exp (None, 48, 48, 96)     1         prune_low_magnitude_block_1_expan
_________________________________________________________________________________________________
prune_low_magnitude_block_1_pad (None, 49, 49, 96)     1         prune_low_magnitude_block_1_expan
_________________________________________________________________________________________________
prune_low_magnitude_block_1_dep (None, 24, 24, 96)     865       prune_low_magnitude_block_1_pad[0
_________________________________________________________________________________________________
prune_low_magnitude_block_1_dep (None, 24, 24, 96)     385       prune_low_magnitude_block_1_depth
_________________________________________________________________________________________________
prune_low_magnitude_block_1_dep (None, 24, 24, 96)     1         prune_low_magnitude_block_1_depth
_________________________________________________________________________________________________
prune_low_magnitude_block_1_pro (None, 24, 24, 24)     4610      prune_low_magnitude_block_1_depth
_________________________________________________________________________________________________
prune_low_magnitude_block_1_pro (None, 24, 24, 24)     97        prune_low_magnitude_block_1_proje
_________________________________________________________________________________________________
prune_low_magnitude_block_2_exp (None, 24, 24, 144)    6914      prune_low_magnitude_block_1_proje
_________________________________________________________________________________________________
prune_low_magnitude_block_2_exp (None, 24, 24, 144)    577       prune_low_magnitude_block_2_expan
_________________________________________________________________________________________________
```

```
prune_low_magnitude_block_2_exp (None, 24, 24, 144)  1          prune_low_magnitude_block_2_expan
prune_low_magnitude_block_2_dep (None, 24, 24, 144)  1297       prune_low_magnitude_block_2_expan
prune_low_magnitude_block_2_dep (None, 24, 24, 144)  577        prune_low_magnitude_block_2_depth
prune_low_magnitude_block_2_dep (None, 24, 24, 144)  1          prune_low_magnitude_block_2_depth
prune_low_magnitude_block_2_pro (None, 24, 24, 24)   6914       prune_low_magnitude_block_2_depth
prune_low_magnitude_block_2_pro (None, 24, 24, 24)   97         prune_low_magnitude_block_2_proje
prune_low_magnitude_block_2_add (None, 24, 24, 24)   1          prune_low_magnitude_block_1_proje
                                                                prune_low_magnitude_block_2_proje
prune_low_magnitude_block_3_exp (None, 24, 24, 144)  6914       prune_low_magnitude_block_2_add[0
prune_low_magnitude_block_3_exp (None, 24, 24, 144)  577        prune_low_magnitude_block_3_expan
prune_low_magnitude_block_3_exp (None, 24, 24, 144)  1          prune_low_magnitude_block_3_expan
prune_low_magnitude_block_3_pad (None, 25, 25, 144)  1          prune_low_magnitude_block_3_expan
prune_low_magnitude_block_3_dep (None, 12, 12, 144)  1297       prune_low_magnitude_block_3_pad[0
prune_low_magnitude_block_3_dep (None, 12, 12, 144)  577        prune_low_magnitude_block_3_depth
prune_low_magnitude_block_3_dep (None, 12, 12, 144)  1          prune_low_magnitude_block_3_depth
prune_low_magnitude_block_3_pro (None, 12, 12, 32)   9218       prune_low_magnitude_block_3_depth
prune_low_magnitude_block_3_pro (None, 12, 12, 32)   129        prune_low_magnitude_block_3_proje
prune_low_magnitude_block_4_exp (None, 12, 12, 192)  12290      prune_low_magnitude_block_3_proje
prune_low_magnitude_block_4_exp (None, 12, 12, 192)  769        prune_low_magnitude_block_4_expan
prune_low_magnitude_block_4_exp (None, 12, 12, 192)  1          prune_low_magnitude_block_4_expan
prune_low_magnitude_block_4_dep (None, 12, 12, 192)  1729       prune_low_magnitude_block_4_expan
prune_low_magnitude_block_4_dep (None, 12, 12, 192)  769        prune_low_magnitude_block_4_depth
prune_low_magnitude_block_4_dep (None, 12, 12, 192)  1          prune_low_magnitude_block_4_depth
prune_low_magnitude_block_4_pro (None, 12, 12, 32)   12290      prune_low_magnitude_block_4_depth
```

prune_low_magnitude_block_4_pro (None, 12, 12, 32)	129	prune_low_magnitude_block_4_proje
prune_low_magnitude_block_4_add (None, 12, 12, 32)	1	prune_low_magnitude_block_3_proje prune_low_magnitude_block_4_proje
prune_low_magnitude_block_5_exp (None, 12, 12, 192)	12290	prune_low_magnitude_block_4_add[0
prune_low_magnitude_block_5_exp (None, 12, 12, 192)	769	prune_low_magnitude_block_5_expan
prune_low_magnitude_block_5_exp (None, 12, 12, 192)	1	prune_low_magnitude_block_5_expan
prune_low_magnitude_block_5_dep (None, 12, 12, 192)	1729	prune_low_magnitude_block_5_expan
prune_low_magnitude_block_5_dep (None, 12, 12, 192)	769	prune_low_magnitude_block_5_depth
prune_low_magnitude_block_5_dep (None, 12, 12, 192)	1	prune_low_magnitude_block_5_depth
prune_low_magnitude_block_5_pro (None, 12, 12, 32)	12290	prune_low_magnitude_block_5_depth
prune_low_magnitude_block_5_pro (None, 12, 12, 32)	129	prune_low_magnitude_block_5_proje
prune_low_magnitude_block_5_add (None, 12, 12, 32)	1	prune_low_magnitude_block_4_add[0 prune_low_magnitude_block_5_proje
prune_low_magnitude_block_6_exp (None, 12, 12, 192)	12290	prune_low_magnitude_block_5_add[0
prune_low_magnitude_block_6_exp (None, 12, 12, 192)	769	prune_low_magnitude_block_6_expan
prune_low_magnitude_block_6_exp (None, 12, 12, 192)	1	prune_low_magnitude_block_6_expan
prune_low_magnitude_block_6_pad (None, 13, 13, 192)	1	prune_low_magnitude_block_6_expan
prune_low_magnitude_block_6_dep (None, 6, 6, 192)	1729	prune_low_magnitude_block_6_pad[0
prune_low_magnitude_block_6_dep (None, 6, 6, 192)	769	prune_low_magnitude_block_6_depth
prune_low_magnitude_block_6_dep (None, 6, 6, 192)	1	prune_low_magnitude_block_6_depth
prune_low_magnitude_block_6_pro (None, 6, 6, 64)	24578	prune_low_magnitude_block_6_depth
prune_low_magnitude_block_6_pro (None, 6, 6, 64)	257	prune_low_magnitude_block_6_proje
prune_low_magnitude_block_7_exp (None, 6, 6, 384)	49154	prune_low_magnitude_block_6_proje
prune_low_magnitude_block_7_exp (None, 6, 6, 384)	1537	prune_low_magnitude_block_7_expan
prune_low_magnitude_block_7_exp (None, 6, 6, 384)	1	prune_low_magnitude_block_7_expan

prune_low_magnitude_block_7_dep (None, 6, 6, 384)	3457	prune_low_magnitude_block_7_expan
prune_low_magnitude_block_7_dep (None, 6, 6, 384)	1537	prune_low_magnitude_block_7_depth
prune_low_magnitude_block_7_dep (None, 6, 6, 384)	1	prune_low_magnitude_block_7_depth
prune_low_magnitude_block_7_pro (None, 6, 6, 64)	49154	prune_low_magnitude_block_7_depth
prune_low_magnitude_block_7_pro (None, 6, 6, 64)	257	prune_low_magnitude_block_7_proje
prune_low_magnitude_block_7_add (None, 6, 6, 64)	1	prune_low_magnitude_block_6_proje prune_low_magnitude_block_7_proje
prune_low_magnitude_block_8_exp (None, 6, 6, 384)	49154	prune_low_magnitude_block_7_add[0
prune_low_magnitude_block_8_exp (None, 6, 6, 384)	1537	prune_low_magnitude_block_8_expan
prune_low_magnitude_block_8_exp (None, 6, 6, 384)	1	prune_low_magnitude_block_8_expan
prune_low_magnitude_block_8_dep (None, 6, 6, 384)	3457	prune_low_magnitude_block_8_expan
prune_low_magnitude_block_8_dep (None, 6, 6, 384)	1537	prune_low_magnitude_block_8_depth
prune_low_magnitude_block_8_dep (None, 6, 6, 384)	1	prune_low_magnitude_block_8_depth
prune_low_magnitude_block_8_pro (None, 6, 6, 64)	49154	prune_low_magnitude_block_8_depth
prune_low_magnitude_block_8_pro (None, 6, 6, 64)	257	prune_low_magnitude_block_8_proje
prune_low_magnitude_block_8_add (None, 6, 6, 64)	1	prune_low_magnitude_block_7_add[0 prune_low_magnitude_block_8_proje
prune_low_magnitude_block_9_exp (None, 6, 6, 384)	49154	prune_low_magnitude_block_8_add[0
prune_low_magnitude_block_9_exp (None, 6, 6, 384)	1537	prune_low_magnitude_block_9_expan
prune_low_magnitude_block_9_exp (None, 6, 6, 384)	1	prune_low_magnitude_block_9_expan
prune_low_magnitude_block_9_dep (None, 6, 6, 384)	3457	prune_low_magnitude_block_9_expan
prune_low_magnitude_block_9_dep (None, 6, 6, 384)	1537	prune_low_magnitude_block_9_depth
prune_low_magnitude_block_9_dep (None, 6, 6, 384)	1	prune_low_magnitude_block_9_depth
prune_low_magnitude_block_9_pro (None, 6, 6, 64)	49154	prune_low_magnitude_block_9_depth
prune_low_magnitude_block_9_pro (None, 6, 6, 64)	257	prune_low_magnitude_block_9_proje

prune_low_magnitude_block_9_add (None, 6, 6, 64)	1	prune_low_magnitude_block_8_add[0 prune_low_magnitude_block_9_proje
prune_low_magnitude_block_10_ex (None, 6, 6, 384)	49154	prune_low_magnitude_block_9_add[0
prune_low_magnitude_block_10_ex (None, 6, 6, 384)	1537	prune_low_magnitude_block_10_expa
prune_low_magnitude_block_10_ex (None, 6, 6, 384)	1	prune_low_magnitude_block_10_expa
prune_low_magnitude_block_10_de (None, 6, 6, 384)	3457	prune_low_magnitude_block_10_expa
prune_low_magnitude_block_10_de (None, 6, 6, 384)	1537	prune_low_magnitude_block_10_dept
prune_low_magnitude_block_10_de (None, 6, 6, 384)	1	prune_low_magnitude_block_10_dept
prune_low_magnitude_block_10_pr (None, 6, 6, 96)	73730	prune_low_magnitude_block_10_dept
prune_low_magnitude_block_10_pr (None, 6, 6, 96)	385	prune_low_magnitude_block_10_proj
prune_low_magnitude_block_11_ex (None, 6, 6, 576)	110594	prune_low_magnitude_block_10_proj
prune_low_magnitude_block_11_ex (None, 6, 6, 576)	2305	prune_low_magnitude_block_11_expa
prune_low_magnitude_block_11_ex (None, 6, 6, 576)	1	prune_low_magnitude_block_11_expa
prune_low_magnitude_block_11_de (None, 6, 6, 576)	5185	prune_low_magnitude_block_11_expa
prune_low_magnitude_block_11_de (None, 6, 6, 576)	2305	prune_low_magnitude_block_11_dept
prune_low_magnitude_block_11_de (None, 6, 6, 576)	1	prune_low_magnitude_block_11_dept
prune_low_magnitude_block_11_pr (None, 6, 6, 96)	110594	prune_low_magnitude_block_11_dept
prune_low_magnitude_block_11_pr (None, 6, 6, 96)	385	prune_low_magnitude_block_11_proj
prune_low_magnitude_block_11_ad (None, 6, 6, 96)	1	prune_low_magnitude_block_10_proj prune_low_magnitude_block_11_proj
prune_low_magnitude_block_12_ex (None, 6, 6, 576)	110594	prune_low_magnitude_block_11_add[
prune_low_magnitude_block_12_ex (None, 6, 6, 576)	2305	prune_low_magnitude_block_12_expa
prune_low_magnitude_block_12_ex (None, 6, 6, 576)	1	prune_low_magnitude_block_12_expa
prune_low_magnitude_block_12_de (None, 6, 6, 576)	5185	prune_low_magnitude_block_12_expa
prune_low_magnitude_block_12_de (None, 6, 6, 576)	2305	prune_low_magnitude_block_12_dept

prune_low_magnitude_block_12_de (None, 6, 6, 576)	1	prune_low_magnitude_block_12_dept
prune_low_magnitude_block_12_pr (None, 6, 6, 96)	110594	prune_low_magnitude_block_12_dept
prune_low_magnitude_block_12_pr (None, 6, 6, 96)	385	prune_low_magnitude_block_12_proj
prune_low_magnitude_block_12_ad (None, 6, 6, 96)	1	prune_low_magnitude_block_11_add[prune_low_magnitude_block_12_proj
prune_low_magnitude_block_13_ex (None, 6, 6, 576)	110594	prune_low_magnitude_block_12_add[
prune_low_magnitude_block_13_ex (None, 6, 6, 576)	2305	prune_low_magnitude_block_13_expa
prune_low_magnitude_block_13_ex (None, 6, 6, 576)	1	prune_low_magnitude_block_13_expa
prune_low_magnitude_block_13_pa (None, 7, 7, 576)	1	prune_low_magnitude_block_13_expa
prune_low_magnitude_block_13_de (None, 3, 3, 576)	5185	prune_low_magnitude_block_13_pad[
prune_low_magnitude_block_13_de (None, 3, 3, 576)	2305	prune_low_magnitude_block_13_dept
prune_low_magnitude_block_13_de (None, 3, 3, 576)	1	prune_low_magnitude_block_13_dept
prune_low_magnitude_block_13_pr (None, 3, 3, 160)	184322	prune_low_magnitude_block_13_dept
prune_low_magnitude_block_13_pr (None, 3, 3, 160)	641	prune_low_magnitude_block_13_proj
prune_low_magnitude_block_14_ex (None, 3, 3, 960)	307202	prune_low_magnitude_block_13_proj
prune_low_magnitude_block_14_ex (None, 3, 3, 960)	3841	prune_low_magnitude_block_14_expa
prune_low_magnitude_block_14_ex (None, 3, 3, 960)	1	prune_low_magnitude_block_14_expa
prune_low_magnitude_block_14_de (None, 3, 3, 960)	8641	prune_low_magnitude_block_14_expa
prune_low_magnitude_block_14_de (None, 3, 3, 960)	3841	prune_low_magnitude_block_14_dept
prune_low_magnitude_block_14_de (None, 3, 3, 960)	1	prune_low_magnitude_block_14_dept
prune_low_magnitude_block_14_pr (None, 3, 3, 160)	307202	prune_low_magnitude_block_14_dept
prune_low_magnitude_block_14_pr (None, 3, 3, 160)	641	prune_low_magnitude_block_14_proj
prune_low_magnitude_block_14_ad (None, 3, 3, 160)	1	prune_low_magnitude_block_13_proj prune_low_magnitude_block_14_proj
prune_low_magnitude_block_15_ex (None, 3, 3, 960)	307202	prune_low_magnitude_block_14_add[

```
prune_low_magnitude_block_15_ex (None, 3, 3, 960)    3841        prune_low_magnitude_block_15_expa
prune_low_magnitude_block_15_ex (None, 3, 3, 960)    1           prune_low_magnitude_block_15_expa
prune_low_magnitude_block_15_de (None, 3, 3, 960)    8641        prune_low_magnitude_block_15_expa
prune_low_magnitude_block_15_de (None, 3, 3, 960)    3841        prune_low_magnitude_block_15_dept
prune_low_magnitude_block_15_de (None, 3, 3, 960)    1           prune_low_magnitude_block_15_dept
prune_low_magnitude_block_15_pr (None, 3, 3, 160)    307202      prune_low_magnitude_block_15_dept
prune_low_magnitude_block_15_pr (None, 3, 3, 160)    641         prune_low_magnitude_block_15_proj
prune_low_magnitude_block_15_ad (None, 3, 3, 160)    1           prune_low_magnitude_block_14_add[
                                                                 prune_low_magnitude_block_15_proj
prune_low_magnitude_block_16_ex (None, 3, 3, 960)    307202      prune_low_magnitude_block_15_add[
prune_low_magnitude_block_16_ex (None, 3, 3, 960)    3841        prune_low_magnitude_block_16_expa
prune_low_magnitude_block_16_ex (None, 3, 3, 960)    1           prune_low_magnitude_block_16_expa
prune_low_magnitude_block_16_de (None, 3, 3, 960)    8641        prune_low_magnitude_block_16_expa
prune_low_magnitude_block_16_de (None, 3, 3, 960)    3841        prune_low_magnitude_block_16_dept
prune_low_magnitude_block_16_de (None, 3, 3, 960)    1           prune_low_magnitude_block_16_dept
prune_low_magnitude_block_16_pr (None, 3, 3, 320)    614402      prune_low_magnitude_block_16_dept
prune_low_magnitude_block_16_pr (None, 3, 3, 320)    1281        prune_low_magnitude_block_16_proj
prune_low_magnitude_Conv_1 (Pru (None, 3, 3, 1280)   819202      prune_low_magnitude_block_16_proj
prune_low_magnitude_Conv_1_bn ( (None, 3, 3, 1280)   5121        prune_low_magnitude_Conv_1[0][0]
prune_low_magnitude_out_relu (P (None, 3, 3, 1280)   1           prune_low_magnitude_Conv_1_bn[0][
prune_low_magnitude_global_aver (None, 1280)         1           prune_low_magnitude_out_relu[0][0
prune_low_magnitude_dense (Prun (None, 1)            2563        prune_low_magnitude_global_averag
==================================================================================================
Total params: 4,386,273
Trainable params: 2,225,153
Non-trainable params: 2,161,120
```

1.2 Pruning log

```
Epoch 1/10
250/250 [==============================] - 47s 186ms/step - loss: 0.1706 - acc: 0.9315 - val_loss: 0.1633 - val_acc: 0.9410
Epoch 2/10
250/250 [==============================] - 46s 183ms/step - loss: 0.1761 - acc: 0.9265 - val_loss: 0.1581 - val_acc: 0.9360
Epoch 3/10
250/250 [==============================] - 46s 184ms/step - loss: 0.1417 - acc: 0.9410 - val_loss: 0.1561 - val_acc: 0.9310
Epoch 4/10
250/250 [==============================] - 47s 188ms/step - loss: 0.1371 - acc: 0.9485 - val_loss: 0.1692 - val_acc: 0.9350
Epoch 5/10
250/250 [==============================] - 47s 187ms/step - loss: 0.1067 - acc: 0.9635 - val_loss: 0.1716 - val_acc: 0.9280
Epoch 6/10
250/250 [==============================] - 47s 187ms/step - loss: 0.1230 - acc: 0.9555 - val_loss: 0.1624 - val_acc: 0.9360
Epoch 7/10
250/250 [==============================] - 47s 187ms/step - loss: 0.1134 - acc: 0.9565 - val_loss: 0.1544 - val_acc: 0.9400
Epoch 8/10
250/250 [==============================] - 47s 187ms/step - loss: 0.1161 - acc: 0.9585 - val_loss: 0.1467 - val_acc: 0.9400
Epoch 9/10
250/250 [==============================] - 47s 187ms/step - loss: 0.0901 - acc: 0.9685 - val_loss: 0.1462 - val_acc: 0.9450
Epoch 10/10
250/250 [==============================] - 47s 188ms/step - loss: 0.0794 - acc: 0.9720 - val_loss: 0.1462 - val_acc: 0.9370
```

1.3 Comparison

Final sparsity	0.3 (original)	0.25	0.5	0.75
Compressed model size (MB)	6.43	6.73	5.11	3.29
Validation acc	94.50%	94.10%	90.20%	63.80%

1.4 Description of the steps of the pruning process

Step 1: Start the training job

Step 2: Acquire the weights, gradients, biases, and activation outputs

Step 3: Compute filter ranks

Step 4: Prune low-ranking filters

Step 5: Set new weights

Step 6: Start the training job with the pruned model

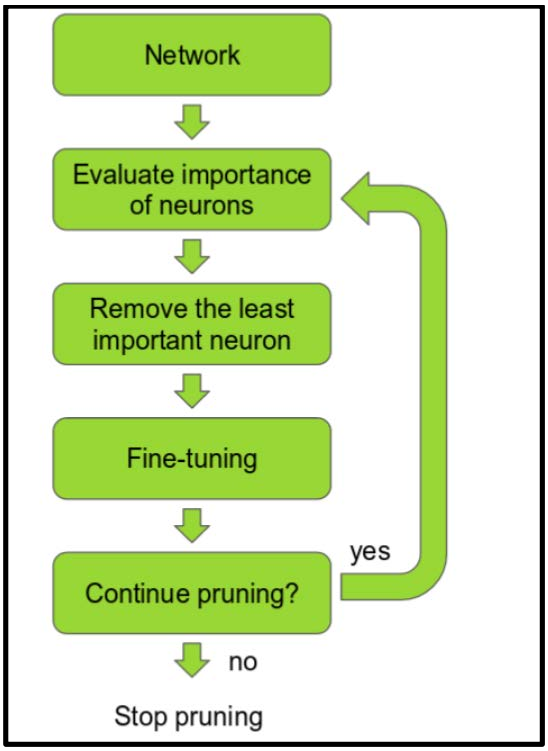

taken from https://jacobgil.github.io/deeplearning/pruning-deep-learning

2. Task 2

2.1 Quantization process

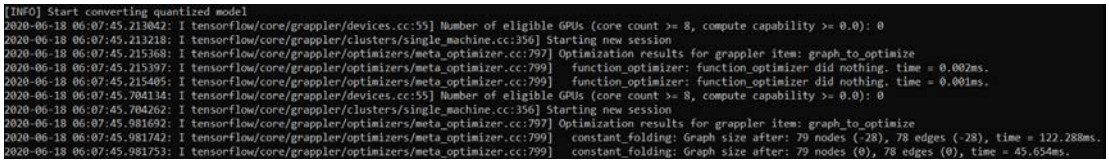

2.2 Validation accuracy

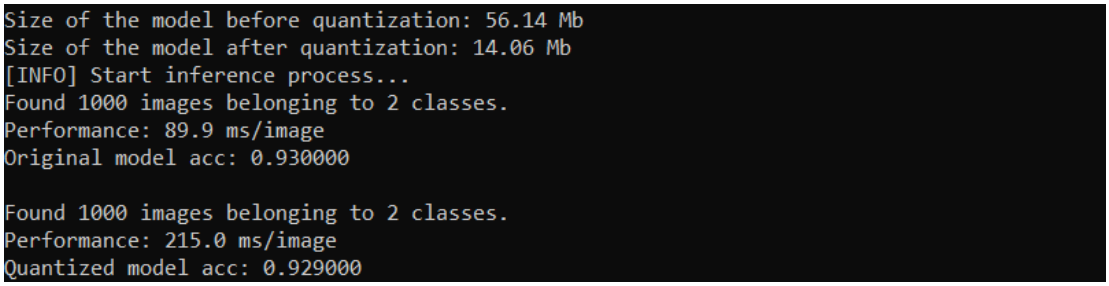

2.3 Comparison

Quantization	None (Original)	Dynamic Range	FP16
tflite file's size (MB)	56.14	14.06	28.08
Validation acc	93.00%	92.90%	93.00%

2.4 Description of the difference between post-training quantization and quantization aware training

There are two forms of quantization:

a. Post-training quantization
   - Easier to use
   - Includes techniques to reduce CPU and hardware accelerator latency, processing power, and model size with little degradation in model accuracy.

b. Quantization aware training
   - Often better for model accuracy
   - Emulates inference-time quantization, creating a model that downstream tools will use to produce the quantized model.
   - The quantized models use lower-precision, such as 8-bit, leading to benefits during deployment.

3. Advance

3.1 Modifications

- The value of epochs
  I changed the value of epochs to 15 for both training on VGG16 model and pruned VGG16 model
- The width and height of images
  I modified the size of images from (112, 112) to (150, 150)
- Final sparsity
  As stated in hw3.pdf, the final sparsity must be set as 0.9
- Callbacks
  I have added a learning reduction to callbacks during the model training

```
# Set a learning rate annealer
learning_rate_reduction = ReduceLROnPlateau(monitor='val_acc',
                                            patience=3,
                                            verbose=1,
                                            factor=0.5,
                                            min_lr=0.00001)
```

During VGG16 training

```
# Create model
model = create_model()
print('[INFO] Start training process...')

model.fit(
      train_generator,
      steps_per_epoch=train_generator.__len__(),
      epochs=EPOCHS,
      validation_data=validation_generator,
      validation_steps=validation_generator.__len__(),
      callbacks=[learning_rate_reduction]
)

model_path = './models/VGG16_model.h5'

print('[INFO] Save model to {}'.format(model_path))
tf.keras.models.save_model(model, model_path, include_optimizer=False)
```

During pruned_VGG16 training

```
# Assign pruning paramaters
pruned_model = sparsity.prune_low_magnitude(model, **pruning_params)

# Print the converted model
pruned_model.summary()

pruned_model.compile(loss='binary_crossentropy', optimizer=OPTIMIZERS, metrics=['acc'])

callbacks = [
    sparsity.UpdatePruningStep(),
    sparsity.PruningSummaries(log_dir='./', profile_batch=0),
    learning_rate_reduction
]

print('[INFO] Start pruning process...')

pruned_model.fit(
      train_generator,
      steps_per_epoch=train_generator.__len__(),
      callbacks=callbacks,
      epochs=epochs,
      validation_data=validation_generator,
      validation_steps=validation_generator.__len__()
)

pruned_model_path = './models/pruned_VGG16.h5'
# convert pruned model to original
final_model = sparsity.strip_pruning(pruned_model)
tf.keras.models.save_model(final_model, pruned_model_path, include_optimizer=False)
```

Learning Reduction during the training process

```
Epoch 8/15
250/250 [==============================] - 55s 218ms/step - loss: 0.3870 - acc: 0.8075 - val_loss: 0.1704 - val_acc: 0.9280
Epoch 9/15
Epoch 00009: ReduceLROnPlateau reducing learning rate to 0.004999999888241291.
Epoch 10/15
250/250 [==============================] - 56s 224ms/step - loss: 0.0078 - acc: 0.9990 - val_loss: 0.1553 - val_acc: 0.9460
Epoch 11/15
```

- Other settings I have tried but haven't added them
  I have tried to modify the optimizers such as Adam, Adagrad, Adamax, Nadam, Ftrl, RMSprop, and so on, but the result comes out the SGD would be the best optimizer on this case.

3.2 Results

- VGG model training
  Model summary:

Input shape changed to (150, 150, 3)

```
Model: "model"
_________________________________________________________________
Layer (type)                 Output Shape              Param #
=================================================================
input_1 (InputLayer)         [(None, 150, 150, 3)]     0
_________________________________________________________________
block1_conv1 (Conv2D)        (None, 150, 150, 64)      1792
_________________________________________________________________
block1_conv2 (Conv2D)        (None, 150, 150, 64)      36928
_________________________________________________________________
block1_pool (MaxPooling2D)   (None, 75, 75, 64)        0
_________________________________________________________________
block2_conv1 (Conv2D)        (None, 75, 75, 128)       73856
_________________________________________________________________
block2_conv2 (Conv2D)        (None, 75, 75, 128)       147584
_________________________________________________________________
block2_pool (MaxPooling2D)   (None, 37, 37, 128)       0
_________________________________________________________________
block3_conv1 (Conv2D)        (None, 37, 37, 256)       295168
_________________________________________________________________
block3_conv2 (Conv2D)        (None, 37, 37, 256)       590080
_________________________________________________________________
block3_conv3 (Conv2D)        (None, 37, 37, 256)       590080
```

```
block3_pool (MaxPooling2D)   (None, 18, 18, 256)       0
_________________________________________________________________
block4_conv1 (Conv2D)        (None, 18, 18, 512)       1180160
_________________________________________________________________
block4_conv2 (Conv2D)        (None, 18, 18, 512)       2359808
_________________________________________________________________
block4_conv3 (Conv2D)        (None, 18, 18, 512)       2359808
_________________________________________________________________
block4_pool (MaxPooling2D)   (None, 9, 9, 512)         0
_________________________________________________________________
block5_conv1 (Conv2D)        (None, 9, 9, 512)         2359808
_________________________________________________________________
block5_conv2 (Conv2D)        (None, 9, 9, 512)         2359808
_________________________________________________________________
block5_conv3 (Conv2D)        (None, 9, 9, 512)         2359808
_________________________________________________________________
block5_pool (MaxPooling2D)   (None, 4, 4, 512)         0
_________________________________________________________________
global_average_pooling2d (Gl (None, 512)               0
_________________________________________________________________
dense (Dense)                (None, 1)                 513
=================================================================
Total params: 14,715,201
Trainable params: 14,715,201
Non-trainable params: 0
_________________________________________________________________
```

Training

```
Epoch 8/15
250/250 [==============================] - 55s 218ms/step - loss: 0.3870 - acc: 0.8075 - val_loss: 0.1704 - val_acc: 0.9280
Epoch 9/15
249/250 [============================>.] - ETA: 0s - loss: 0.0567 - acc: 0.9809
Epoch 00009: ReduceLROnPlateau reducing learning rate to 0.004999999888241291.
250/250 [==============================] - 55s 218ms/step - loss: 0.0566 - acc: 0.9810 - val_loss: 0.1544 - val_acc: 0.9450
Epoch 10/15
250/250 [==============================] - 56s 224ms/step - loss: 0.0078 - acc: 0.9990 - val_loss: 0.1553 - val_acc: 0.9460
Epoch 11/15
250/250 [==============================] - 55s 221ms/step - loss: 0.0017 - acc: 1.0000 - val_loss: 0.1855 - val_acc: 0.9440
Epoch 12/15
249/250 [============================>.] - ETA: 0s - loss: 9.2023e-04 - acc: 1.0000
Epoch 00012: ReduceLROnPlateau reducing learning rate to 0.0024999999441206455.
250/250 [==============================] - 56s 223ms/step - loss: 9.1658e-04 - acc: 1.0000 - val_loss: 0.1848 - val_acc: 0.9500
Epoch 13/15
250/250 [==============================] - 55s 221ms/step - loss: 6.3105e-04 - acc: 1.0000 - val_loss: 0.1925 - val_acc: 0.9470
Epoch 14/15
250/250 [==============================] - 55s 221ms/step - loss: 5.3363e-04 - acc: 1.0000 - val_loss: 0.1956 - val_acc: 0.9490
Epoch 15/15
249/250 [============================>.] - ETA: 0s - loss: 4.6569e-04 - acc: 1.0000
Epoch 00015: ReduceLROnPlateau reducing learning rate to 0.0012499999720603228.
250/250 [==============================] - 55s 221ms/step - loss: 4.6403e-04 - acc: 1.0000 - val_loss: 0.1985 - val_acc: 0.9500
[INFO] Save model to ./models/VGG16_model.h5
```

- Pruning

  Model summary

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 150, 150, 3)]	0
prune_low_magnitude_block1_c	(None, 150, 150, 64)	3522
prune_low_magnitude_block1_c	(None, 150, 150, 64)	73794
prune_low_magnitude_block1_p	(None, 75, 75, 64)	1
prune_low_magnitude_block2_c	(None, 75, 75, 128)	147586
prune_low_magnitude_block2_c	(None, 75, 75, 128)	295042
prune_low_magnitude_block2_p	(None, 37, 37, 128)	1
prune_low_magnitude_block3_c	(None, 37, 37, 256)	590082
prune_low_magnitude_block3_c	(None, 37, 37, 256)	1179906
prune_low_magnitude_block3_c	(None, 37, 37, 256)	1179906
prune_low_magnitude_block3_p	(None, 18, 18, 256)	1
prune_low_magnitude_block4_c	(None, 18, 18, 512)	2359810
prune_low_magnitude_block4_c	(None, 18, 18, 512)	4719106
prune_low_magnitude_block4_c	(None, 18, 18, 512)	4719106
prune_low_magnitude_block4_p	(None, 9, 9, 512)	1
prune_low_magnitude_block5_c	(None, 9, 9, 512)	4719106
prune_low_magnitude_block5_c	(None, 9, 9, 512)	4719106
prune_low_magnitude_block5_c	(None, 9, 9, 512)	4719106
prune_low_magnitude_block5_p	(None, 4, 4, 512)	1
prune_low_magnitude_global_a	(None, 512)	1
prune_low_magnitude_dense (P	(None, 1)	1027

Total params: 29,426,211
Trainable params: 14,715,201
Non-trainable params: 14,711,010

Pruning and model recovering

Training

```
Epoch 00006: ReduceLROnPlateau reducing learning rate to 0.0049999999888241291.
250/250 [==============================] - 60s 239ms/step - loss: 0.0271 - acc: 0.9900 - val_loss: 0.2046 - val_acc: 0.9250
Epoch 7/15
250/250 [==============================] - 59s 238ms/step - loss: 0.0660 - acc: 0.9830 - val_loss: 0.1757 - val_acc: 0.9310
Epoch 8/15
250/250 [==============================] - 60s 239ms/step - loss: 0.0517 - acc: 0.9850 - val_loss: 0.2142 - val_acc: 0.9250
Epoch 9/15
249/250 [============================>.] - ETA: 0s - loss: 0.2851 - acc: 0.8830
Epoch 00009: ReduceLROnPlateau reducing learning rate to 0.0024999999441206455.
250/250 [==============================] - 60s 239ms/step - loss: 0.2853 - acc: 0.8830 - val_loss: 0.3198 - val_acc: 0.8440
Epoch 10/15
250/250 [==============================] - 60s 239ms/step - loss: 0.1623 - acc: 0.9460 - val_loss: 0.2478 - val_acc: 0.8950
Epoch 11/15
250/250 [==============================] - 60s 238ms/step - loss: 0.2252 - acc: 0.9115 - val_loss: 0.2465 - val_acc: 0.9030
Epoch 12/15
249/250 [============================>.] - ETA: 0s - loss: 0.2181 - acc: 0.9217
Epoch 00012: ReduceLROnPlateau reducing learning rate to 0.0012499999720603228.
250/250 [==============================] - 60s 239ms/step - loss: 0.2176 - acc: 0.9220 - val_loss: 0.2478 - val_acc: 0.8960
Epoch 13/15
250/250 [==============================] - 59s 238ms/step - loss: 0.2076 - acc: 0.9290 - val_loss: 0.3046 - val_acc: 0.8660
Epoch 14/15
250/250 [==============================] - 60s 239ms/step - loss: 0.1644 - acc: 0.9410 - val_loss: 0.2522 - val_acc: 0.8960
Epoch 15/15
249/250 [============================>.] - ETA: 0s - loss: 0.1408 - acc: 0.9553
Epoch 00015: ReduceLROnPlateau reducing learning rate to 0.0006249999860301614.
250/250 [==============================] - 60s 238ms/step - loss: 0.1407 - acc: 0.9555 - val_loss: 0.2249 - val_acc: 0.9010
```

Comparision of compressed model size and validation accuracy

```
Size of the model before compression: 56.20 MB
Size of the model after compression: 52.20 MB
Size of the pruned model before compression: 56.20 MB
Size of the pruned model after compression: 10.76 MB
[INFO] model val_acc: 0.9010000228881836
[INFO] pruend model val_acc: 0.9010000228881836
```

This is wrong, the actual accuracy achieved by original model is 0.9302123

- Quantization

```
Size of the model before quantization: 56.14 Mb
Size of the model after quantization: 14.05 Mb
```

```
[INFO] Start inference process...
Found 1000 images belonging to 2 classes.
Performance: 122.4 ms/image
Original model acc: 0.901000

Found 1000 images belonging to 2 classes.
Performance: 1048.5 ms/image
Quantized model acc: 0.898000
```